

Memory Management Unit

The Memory Management Unit (MMU) SV32 module is a crucial component in the RISC-V-based processor, serving as the backbone for virtual memory management and address translation.




Figure 1: Input and Outputs of CVA6 MMU SV32

At its core, the MMU SV32 plays a pivotal role in translating virtual addresses into their corresponding physical counterparts. This translation process is paramount for providing memory protection, isolation, and efficient memory management in modern computer systems. Importantly, it handles both instruction and data accesses, ensuring a seamless interaction between the processor and virtual memory.




Figure 2: Major Blocks in CVA6 MMU SV32

The MMU SV32 manages privilege levels and access control, enforcing permissions for user and supervisor modes while handling access exceptions. It employs Translation Lookaside Buffers (TLBs) for efficient address translation, reducing the need for page table access. TLB hits yield quick translations, but on misses, the shared TLB is consulted, and if necessary, the Page Table Walker (PTW) performs page table walks, updating TLBs and managing exceptions during the process.

In addition to these functionalities, the MMU SV32 seamlessly integrates support for Physical Memory Protection (PMP), enabling it to enforce access permissions and memory protection configurations as specified by the PMP settings. This additional layer of security and control enhances the management of memory accesses

1.1 Instruction and Data Interfaces

The MMU SV32 maintains interfaces with the instruction cache (ICache) and the load-store unit (LSU). It receives virtual addresses from these components and proceeds to translate them into physical addresses, a fundamental task for ensuring proper program execution and memory access.

1.2 Signal Description of MMU

Signal	IO	Connection	Type	Description
clk_i	in	Subsystem	logic	Subsystem Clock
rst_ni	in	Subsystem	logic	Asynchronous reset active low
flush_i	in	Controller	logic	Sfence Committed
enable_translation_i	in	CSR RegFile	logic	Indicate address translation request for instruction
en_ld_st_translation_i	in	CSR RegFile	logic	Indicate address translation request for load or store
icache_areq_i	in	Cache Subsystem	icache_arsp_t	Icache Response
icache_areq_o	out	Cache Subsystem	icache_areq_t	Icache Request
misaligned_ex_i	in	Load Store Unit	exception_t	Indicate misaligned exception
lsu_req_i	in	Load Store Unit	logic	Request address translation
lsu_vaddr_i	in	Load Store Unit	logic [riscv::VLEN-1:0]	Virtual Address In
lsu_is_store_i	in	Store Unit	logic	Translation is requested by a store
lsu_dtlb_hit_o	out	Store / Load Unit	logic	Indicate a DTLB hit
lsu_dtlb_ppn_o	out	Load Unit	logic	Send PNN to LSU

			[riscv::PPNW-1:0]	
lsu_valid_o	out	Load Store Unit	logic	Indicate a valid translation
lsu_paddr_o	out	Store / Load Unit	logic [riscv::PLEN-1:0]	Translated Address
lsu_exception_o	out	Store / Load Unit	exception_t	Address Translation threw an exception
priv_lvl_i	in	CSR RegFile	riscv::priv_lvl_t	Privilege level for instruction fetch interface
ld_st_priv_lvl_i	in	CSR RegFile	riscv::priv_lvl_t	Privilege Level for Data Interface
sum_i	in	CSR RegFile	logic	Supervisor User memory access bit in xSTATUS register
mrx_i	in	CSR RegFile	logic	Make Executable Readable bit in xSTATUS CSR register
satp_ppn_l	in	CSR RegFile	logic [riscv::PPNW-1:0]	PPN of top level page table from SATP register
asid_i	in	CSR RegFile	logic [ASID_WIDTH-1:0]	ASID to for the lookup
asid_to_be_flushed	in	Execute Stage	logic [ASID_WIDTH-1:0]	ASID of the entry to be flushed.
vaddr_to_be_flushed_i	in	Execute Stage	logic [riscv::VLEN-1:0]	Virtual address of the entry to be flushed.
flush_tlb_i	in	Controller	logic	SFENCE.VMA committed
itlb_miss_o	out	Performance Counter	logic	Indicate a ITLB miss
dtlb_miss_o	out	Performance Counter	logic	Indicate a DTLB miss
req_port_i	in	Cache Subsystem	dcache_req_o_t	D Cache Data Requests
req_port_o	out	Cache Subsystem	dcache_req_i_t	D Cache Data Response
pmpcfg_i	in	CSR RegFile	riscv::pmpcfg_t [15:0]	PMP configurations
pmpaddr_i	in	CSR RegFile	logic [15:0][riscv::PLEN-3:0]	PMP Address

1.3 Struct Description

icache_areq_t

Signal	Type	Description
fetch_valid	logic	Address Translation Valid
fetch_paddr	logic [riscv::PLEN-1:0]	Physical Address In
fetch_exception	exception_t	Exception occurred during fetch

icache_arsp_t

Signal	Type	Description
fetch_req	logic	Address Translation Request
fetch_vaddr	logic [riscv::VLEN-1:0]	Virtual Address out

exception_t

Signal	Type	Description
cause	riscv::xlen_t	Cause of exception
tval	riscv::xlen_t	Additional information of causing exception (e.g. instruction causing it), address of LD/ST fault
valid	logic	Indicate that exception is valid

pmpcfg_t

Signal	Type	Description
locked	logic	Lock this configuration
reserved	logic[1:0]	Reserved bits in pmpcfg CSR
addr_mode	pmp_addr_mode_t	Addressing Modes: OFF, TOR, NA4, NAPOT
access_type	pmpcfg_access_t	None, read, write, execute

1.4 Control Flow in MMU SV32 Module

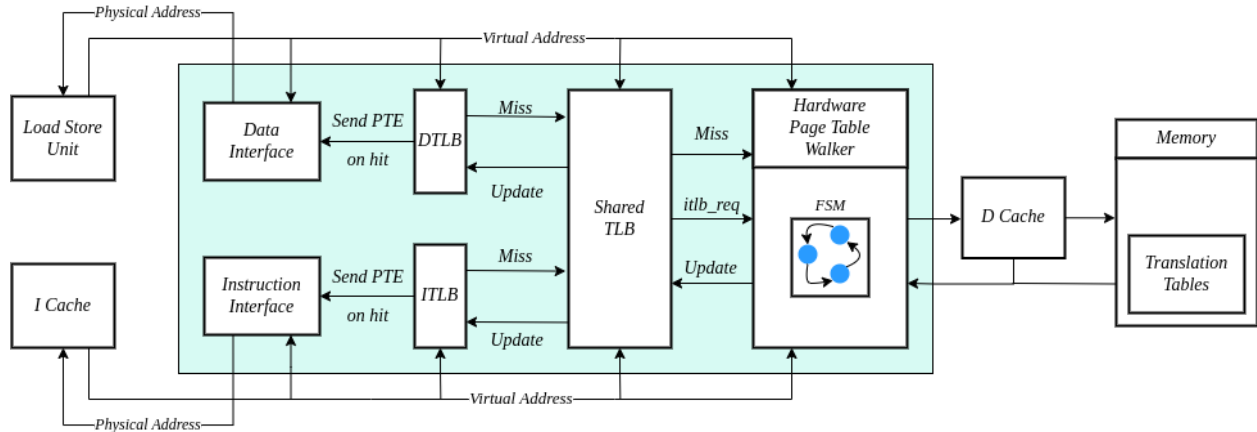


Figure 3: Control Flow in CVA6 MMU SV32

Exception Sources with Address Translation Enabled

Two potential exception sources exist:

- Hardware Page Table Walker (HPTW) throwing an exception, signifying a page fault exception.
- Access error due to insufficient permissions of PMP, known as an access exception.

1.5 Instruction Fetch Interface

The IF stage initiates a request to retrieve memory content at a specific virtual address. When the MMU is disabled, the instruction fetch request is directly passed to the I\$ without modifications.

1.5.1 Address Translation in Instruction Interface

If virtual memory translation is enabled for instruction fetches, the following operations are performed in the instruction interface:

- Compatibility of requested virtual address with selected page based address translation scheme is checked.
- For 4K page translation, the module determines the fetch physical address by combining the physical page number (PPN) from ITLB content and the offset from the virtual address.
- In the case of Mega page translation, if the ITLB indicates a 4M page, the VPN0 from the fetch virtual address is written to the PPN0 of the fetch physical address to ensure alignment for superpage translation.

- If the Instruction TLB (ITLB) lookup hits, the fetch valid signal (which indicates a valid physical address) is activated in response to the input fetch request. Memory region accessibility is checked from the perspective of the fetch operation, potentially triggering a page fault exception in case of an access error or insufficient PMP permission.
- In case of an ITLB miss, if the page table walker (PTW) is active (only active if there is a shared TLB miss) and handling instruction fetches, the fetch valid signal is determined based on PTW errors or access exceptions.

If the fetch physical address doesn't match any execute region, an Instruction Access Fault is raised. When not translating, PMPs are immediately checked against the physical address for access verification.

1.6 Data Interface

Similar to the instruction fetch interface, if the physical address length (PLEN) exceeds the virtual address length (VLEN), the fetch physical address is set by appending zeros to the most significant bits of the virtual address. Simultaneously, the Physical Page Number (PPN) is sent to the Load-Store Unit (LSU) after removing the offset from the virtual address and adding zeros to its most significant bits.

When PLEN is not greater than VLEN, the fetch physical address is truncated to match PLEN, and the PPN is sent to LSU in the same manner.

The address translation process in the data interface is not combinational. Either the load or the store unit will ask for the address translation.

1.6.1 Address Translation in Data Interface

If address translation is enabled for load or store, and no misaligned exception has occurred, the following operations are performed in the data interface:

- Initially, translation is assumed to be invalid, signified by the MMU to LSU.
- The translated physical address is formed by combining the PPN from the Page Table Entry (PTE) and the offset from the virtual address requiring translation. This send one cycle later due to the additional bank of registers which delayed the MMU's answer. The PPN from the PTE is also shared separately with LSU in the same cycle as the hit.
- In the case of superpage translation, as in SV32, known as the 4M page, PPN0 of the translated physical address and the separately shared PPN are updated with the VPN0 of the virtual address.

If a Data TLB (DTLB) hit occurs, it indicates a valid translation, and various fault checks are performed depending on whether it's a load or store request.

- For store requests, if the page is not writable, the dirty flag isn't set, or privileges are violated, it results in a page fault corresponding to the store access. If PMPs are also violated, it leads to an access fault corresponding to the store access. Page faults take precedence over access faults.
- For load requests, a page fault is triggered if there are insufficient access privileges. PMPs are checked again during load access, resulting in an access fault corresponding to load access if PMPs are violated.

In case of a DTLB miss, potential exceptions are monitored during the page table walk. If the PTW indicates a page fault, the corresponding page fault related to the requested type is signaled. If the PTW indicates an access exception, the load access fault is indicated through address translation because the page table walker can only throw load access faults.

1.6.2 Address Translation is Disabled

When address translation is not enabled, the physical address is immediately checked against Physical Memory Protections (PMPs). If there is a request from LSU, no misaligned exception, and PMPs are violated, it results in an access fault corresponding to the request being indicated.

CVA6 Translation Lookaside Buffer SV32

Page tables are accessed for translating virtual memory addresses to physical memory addresses. This translation needs to be carried out for every load and store instruction and also for every instruction fetch. Since page tables are resident in physical memory, accessing these tables in all these situations has a significant impact on performance. Page table accesses occur in patterns that are closely related in time. Furthermore, the spatial and temporal locality of data accesses or instruction fetches mean that the same page is referenced repeatedly. Taking advantage of these access patterns the processor keeps the information of recent address translations, to enable fast retrieval, in a small cache called the Translation Lookaside Buffer (TLB) or an address-translation cache.

The CVA6 TLB is structured as a fully associative cache, where the virtual address that needs to be translated is compared against all the individual TLB entries. Given a virtual address, the processor examines the TLB (*TLB lookup*) to determine if the virtual page number (VPN) of the page being accessed is in the TLB. When a TLB entry is found (*TLB hit*), the TLB returns the corresponding physical page number (PPN) which is used to calculate the target physical address. If no TLB entry is found (*TLB miss*) the

processor has to read individual page table entries from memory (*Table walk*). In CVA6 table walking is supported by dedicated hardware. Once the processor finishes the table walk it has the Physical Page Number (*PPN*) corresponding to the Virtual Page Number (*VPN*) That needs to be translated. The processor adds an entry for this address translation to the TLB so future translations of that virtual address will happen quickly through the TLB. During the table walk the processor may find out that the corresponding physical page is not resident in memory. At this stage a page table exception (*Page Fault*) is generated which gets handled by the operating system. The operating system places the appropriate page in memory, updates the appropriate page tables and returns execution to the instruction which generated the exception.

The inputs and output signals of the TLB are shown in the following two figures.

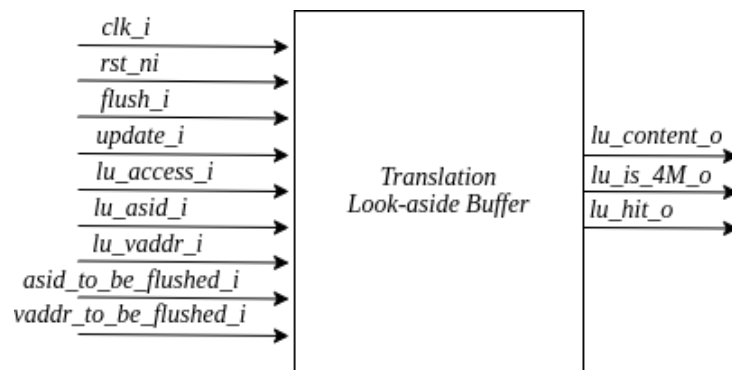


Figure 4: Inputs and Outputs of CVA6 TLB

2.1 Signal Description

Signal	IO	Connection	Type	Description
clk_i	in	Subsystem	logic	Subsystem Clock
rst_ni	in	Subsystem	logic	Asynchronous reset active low
flush_i	in	Controller	logic	TLB flush request
update_i	in	Shared TLB	tlb_update_sv32_t	Updated tag and content of TLB.
lu_access_i	in	Cache Subsystem	logic	Signal indicating a lookup access is being requested.
lu_asid_i	in	CSR RegFile	Logic [ASID_WIDTH-1:0]	ASID (Address Space Identifier) for the lookup.
lu_vaddr_i	in	Cache Subsystem	logic [riscv::VLEN-1:0]	Virtual address for the lookup.

lu_content_o	out	MMU SV32	riscv::pte_sv32_t	Output for the content of the TLB entry.
asid_to_be_flushed_i	in	Execute Stage	Logic [ASID_WIDTH-1:0]	ASID of the entry to be flushed.
vaddr_to_be_flushed_i	in	Execute Stage	Logic [riscv::VLEN-1:0]	Virtual address of the entry to be flushed.
lu_is_4M_o	out	MMU SV32	logic	Output indicating whether the TLB entry corresponds to a 4MB page.
lu_hit_o	out	MMU SV32	logic	Output indicating whether the lookup resulted in a hit or miss

2.2 Struct Description

tlb_update_sv32_t

Signal	Type	Description
valid	logic	Indicates whether the TLB entry is valid or not 0: TLB entry is not valid 1: TLB entry is valid
is_4M	logic	Indicates if the TLB entry corresponds to a 4MB page.
vpn	logic[19:0]	Virtual Page Number (VPN) used for updating the TLB, consisting of 20 bits.
asid	logic [8:0]	Address Space Identifier (ASID) used for updating the TLB, with a length of 9 bits for Sv32 MMU.
content	pte_sv32_t	Content of the TLB update entry, defined by the structure

pte_sv32_t

Signal	Type	Description
ppn	logic [21 :0]	22 bit Physical Page Number (PPN)
rsw	logic[1:0]	Reserved for use by supervisor software
d	logic	Dirty bit indicating whether the page has been modified (dirty) or not 0: Page is clean i.e., has not been written 1: Page is dirty i.e., has been written
a	logic	Accessed bit indicating whether the page has been accessed 0: Virtual page has not been accessed since the last time A bit was cleared 1: Virtual page has been read, written, or fetched from since the

		last time the A bit was cleared
g	logic	Global bit marking a page as part of a global address space valid for all ASIDs 0: Translation is valid for specific ASID 1: Translation is valid for all ASIDs
u	logic	User bit indicating privilege level of the page 0: Page is not accessible in user mode but in supervisor mode 1: Page is accessible in user mode but not in supervisor mode.
x	logic	Execute bit which allows execution of code from the page 0: Code execution is not allowed 1: Code execution is permitted
w	logic	Write bit allows the page to be written 0: Write operations are not allowed 1: Write operations are permitted
r	logic	Read bit allows read access to the page. 0: Read operations are not allowed 1: Read operations are permitted
v	logic	Indicates if the entry in memory is valid or not 0: Page is invalid i.e. page is not in physical memory 1: Page is valid i.e. page resides in physical memory

2.3 TLB Entry Fields

The number of TLB entries can be changed via a design parameter. In 32-bit configurations of CVA6 only 2 TLB entries are instantiated. Each TLB entry is made up of two fields: Tag and Content. The Tag field holds the virtual page number (VPN1, VPN0), ASID, page size (is_4M) along with a valid bit (VALID) indicating that the entry is valid. The SV32 virtual page number, which is supported by CV32A6X, is further split into two separate virtual page numbers VPN1 and VPN0. The Content field contains two physical page numbers (PPN1, PPN0) along with a number of bits which specify various attributes of the physical page. Note that the V bit in the Content field is the V bit which is present in the page table in memory. It is copied from the page table, as is, and the VALID bit in the Tag is set based on its value. The TLB entry fields are shown in **Figure 5**.

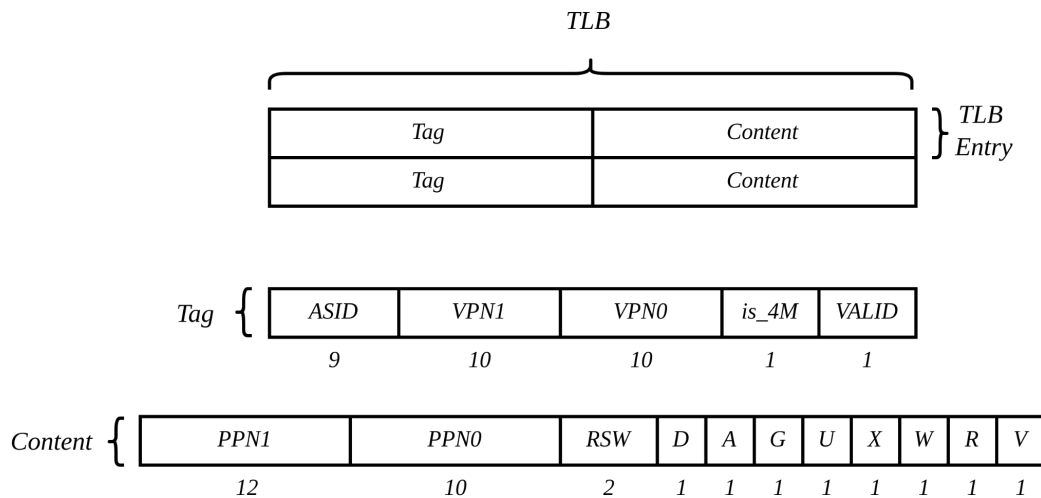


Figure 5: Fields in CVA6 TLB entry

2.4 CVA6 TLB Management / Implementation

The CVA6 TLB implements the following three functions:

- **Translation:** This function implements the address lookup and match logic
- **Update and Flush:** This function implements the update and flush logic
- **Pseudo Least Recently Used Replacement Policy:** This function implements the replacement policy for TLB entries.

2.4.1 Translation

This function takes in the virtual address and certain other fields, examines the TLB to determine if the virtual page number of the page being accessed is in the TLB or not. If a TLB entry is found (TLB hit), the TLB returns the corresponding physical page number (PPN) which is then used to calculate the target physical address. The following checks are done as part of this lookup function to find a match in the TLB:

- **Validity Check:** For a TLB hit, the associated TLB entry must be valid .
- **ASID and Global Flag Check:** The TLB entry's ASID must match the given ASID (ASID associated with the Virtual address). If the TLB entry's Global bit (G) bit is set then this check is not done. This ensures that the translation is either specific to the provided ASID or it is globally applicable.

- **Level 1 VPN match:** SV32 implements a two-level page table. As such the virtual address is broken up into three parts which are the virtual page number 1, virtual page number 0 and offset. So the condition that is checked next is that the virtual page number 1 of the virtual address matches the virtual page number 1 (VPN1) of the TLB entry.
- **Level 0 VPN match or 4-Mega Page:** The last condition to be checked, for a TLB hit, is that the virtual page number 0 of the virtual address matches the virtual page number 0 of the TLB entry (VPN0). This match is ignored if the `is_4M` bit in the Tag is set which implies a super 4M page.

All the conditions listed above are checked against every TLB entry. If there is a TLB hit then the corresponding bit in the hit array is set. **Figure 6** Illustrates the TLB hit/miss process listed above.

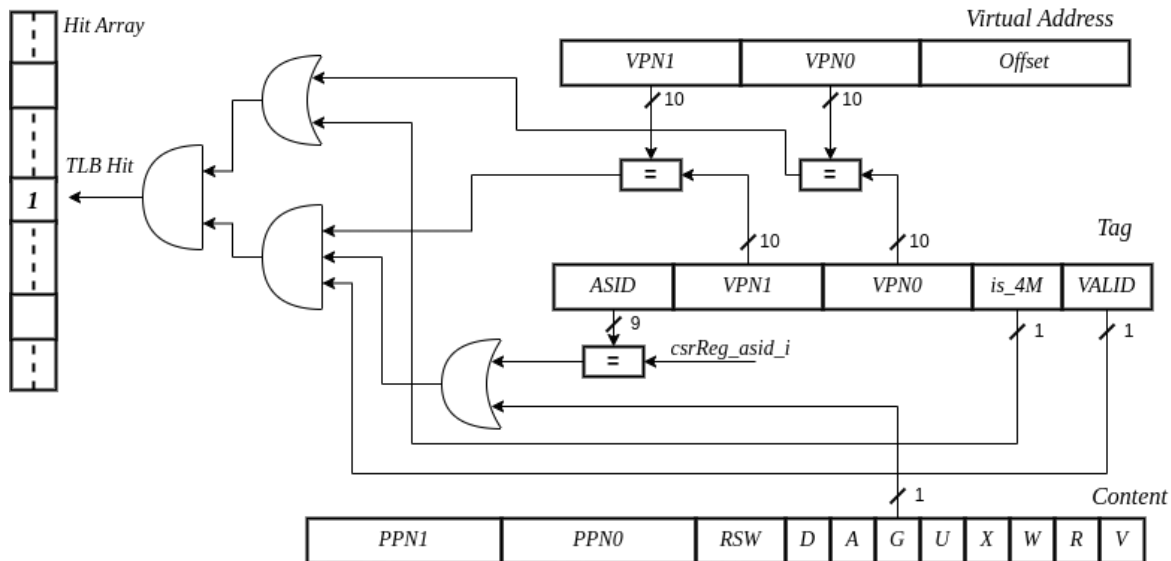


Figure 6: Block diagram of CVA6 TLB hit or miss

2.4.2 Flushing TLB entries

The `SFENCE.VMA` instruction can be used with certain specific source register specifiers (`rs1` & `rs2`) to flush a specific TLB entry, some set of TLB entries or all TLB entries. Like all instructions this action only takes place when the `SFENCE.VMA` instruction is committed (shown via the `commit_sfence` signal in the following figures.) The behavior of the instruction is as follows:

- **If `rs1` is not equal to `x0` and `rs2` is not equal to `x0`:** Invalidate all TLB entries which contain leaf page table entries corresponding to the virtual address in `rs1` (shown below as *Virtual Address to be flushed*) and that match the address

space identifier as specified by integer register rs2 (shown below as *asid_to_be_flushed_i*), except for entries containing global mappings. This is referred to as the “*SFENCE.VMA vaddr asid*” case.

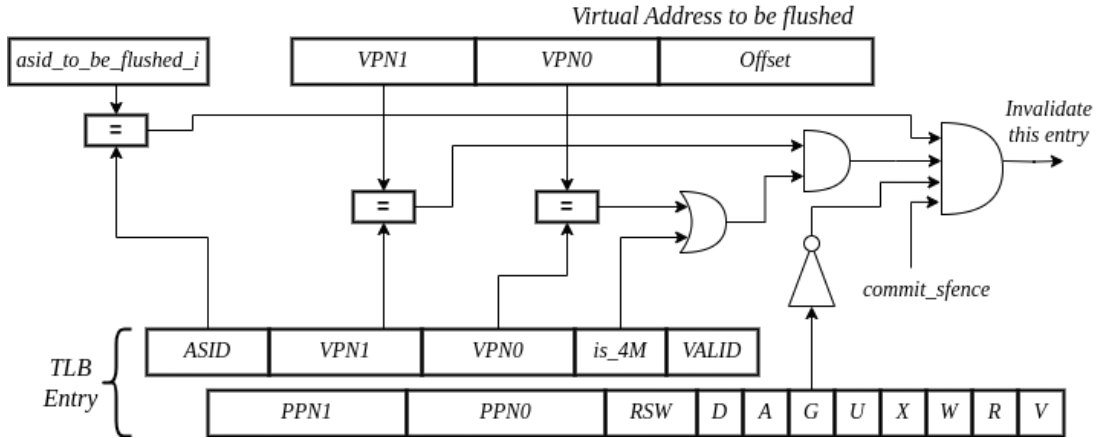


Figure 7: Invalidate TLB entry if ASID and virtual address match

- **If rs1 is equal to x0 and rs2 is equal to x0:** Invalidate all TLB entries for all address spaces. This is referred to as the "*SFENCE.VMA x0 x0*" case.

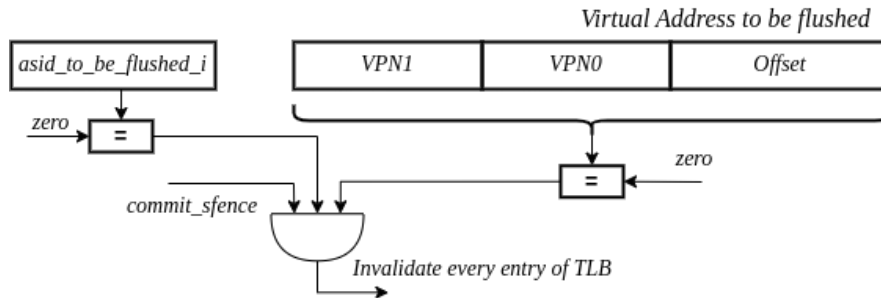


Figure 8: Invalidate all TLB entries if both source register specifiers are x0

- **If rs1 is not equal to x0 and rs2 is equal to x0:** invalidate all TLB entries that contain leaf page table entries corresponding to the virtual address in rs1, for all address spaces. This is referred to as the “*SFENCE.VMA vaddr x0*” case.

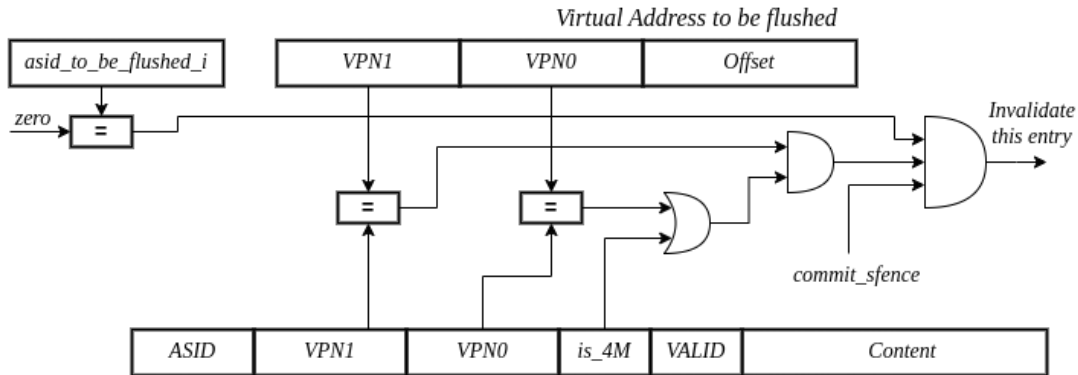


Figure 9: Invalidate TLB entry with matching virtual address for all address spaces

- **If `rs1` is equal to `x0` and `rs2` is not equal to `x0`:** Invalidate all TLB entries matching the address space identified by integer register `rs2`, except for entries containing global mappings. This is referred to as the “*SFENCE.VMA 0 asid*” case.

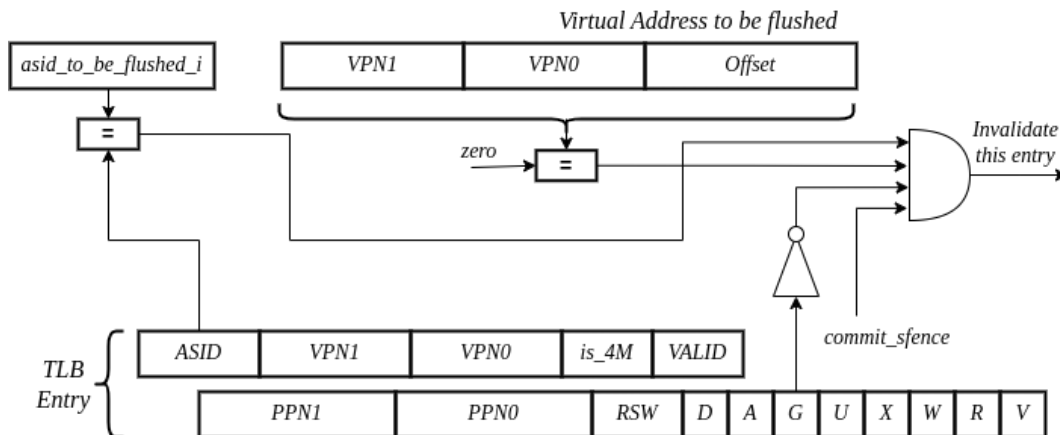


Figure 10: Invalidate TLB entry for matching ASIDs

2.4.3 Updating TLB

When a TLB valid update request is signaled by the shared TLB, and the replacement policy select the update of a specific TLB entry, the corresponding entry's tag is updated with the new tag, and its associated content is refreshed with the information from the update request. This ensures that the TLB entry accurately reflects the new translation information.

2.4.4 Pseudo Least Recently Used Replacement (PLRU) Policy

Cache replacement algorithms are used to determine which TLB entry should be replaced, because it is not likely to be used in the near future. The Pseudo-Least-Recently-Used (PLRU) is a cache entry replacement algorithm, derived from Least-Recently-Used (LRU) cache entry replacement algorithm, used by the TLB. Instead of precisely tracking recent usage as the LRU algorithm does, PLRU employs an approximate measure to determine which entry in the cache has not been recently used and as such can be replaced.

CVA6 implements the PLRU algorithm via the Tree-PLRU method which implements a binary tree. The TLB entries are the leaf nodes of the tree. Each internal node, of the tree, consists of a single bit, referred to as the state bit or *plru bit*, indicating which subtree contains the (pseudo) least recently used entry (the PLRU); 0 for the left hand tree and 1 for the right hand tree. Following this traversal, the leaf node reached, corresponds to the PLRU entry which can be replaced. Having accessed an entry (so as to replace it) we need to promote that entry to be the Most Recently Used (MRU) entry. This is done by updating the value of each node along the access path to point away from that entry. If the accessed entry is a right child i.e., its parent node value is 1, it is set to 0, and if the parent is the left child of its parent (the grandparent of the accessed node) then its node value is set to 1 and so on all the way up to the root node.

The PLRU binary tree is implemented as an array of node values. Nodes are organized in the array based on levels, with those from lower levels appearing before higher ones. Furthermore those on the left side of a node appear before those on the right side of a node. The figure below shows a tree and the corresponding array.




Figure 11: PLRU Tree Indexing

Traversing the tree by following the direction as listed in the nodes at every level gets us to the PLRU entry which is D. Traversing the tree by *not* following the direction as listed in the nodes at every level gets us to the MRU entry which is F.

For a n-way associative TLB we require $n - 1$ internal nodes in the tree. With those nodes, two operations need to be performed efficiently.

- Promote the accessed entry to be MRU
- Identify which entry to replace (i.e. the PLRU entry)

2.4.4.1 Updating the PLRU-Tree

For a TLB entry which is accessed, the following steps are taken to make it the MRU:

1. Iterate through each level of the binary tree.
2. Calculate the index of the leftmost child within the current level. Let us call that index the index base.
3. Calculate the shift amount to identify the relevant node based on the level and TLB entry index.
4. Calculate the new value that the node should have in order to make the accessed entry the Most Recently Used (MRU). The new value of the root node is the opposite of the TLB entry index, MSB at the root node, MSB - 1 at the next level and so on.

- Assign this new value to the relevant node, ensuring that the hit entry becomes the MRU within the binary tree structure.

At level 0, no bit of the TLB entry's index determines the offset from the index base because it's a root node. At level 1, MSB of entry's index determines the amount of offset from index base at that level. At level 2, the first two bits of the entry's index from MSB side determine the offset from the index base because there are 4 nodes at the level 2 and so on.

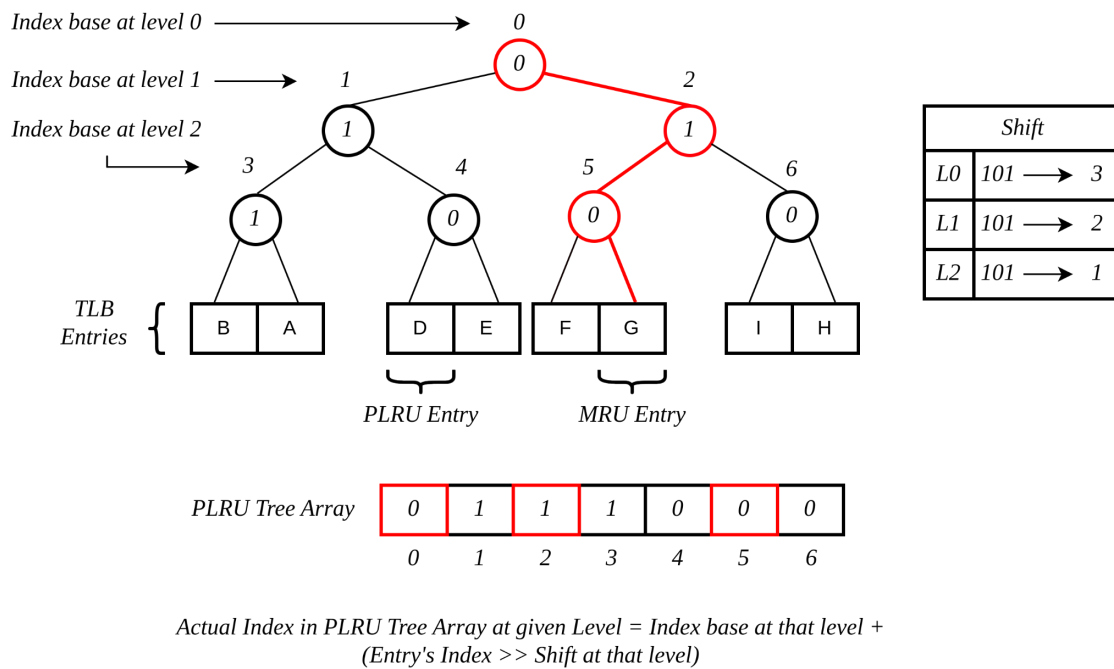


Figure 12: Promote Entry to be MRU

In the above figure entry at index 5, is accessed. To make it MRU entry, every node along the access path should point away from it. Entry 5 is a right child, therefore, its parent plru bit set to 0, its parent is a left child, its grand parent's plru bit set to 1, and great grandparent's plru bit set to 0.

2.4.4.2 Entry Selection for Replacement

Every TLB entry is checked for the replacement entry. The following steps are taken:

- Iterate through each level of the binary tree.

2. Calculate the index of the leftmost child within the current level. Let us call that index the index base.
3. Calculate the shift amount to identify the relevant node based on the level and TLB entry index.
4. If the corresponding bit of the entry's index matches the value of the node being traversed at the current level, keep the replacement signal high for that entry; otherwise, set the replacement signal to low.




Figure 13: Possible path traverse for entry selection for replacement

Figure shows every possible path that traverses to find out the PLRU entry. If the plru bit at each level matches with the corresponding bit of the entry's index, that's the next entry to replace. Below Table shows the entry selection for replacement.

Path Traverse	PLRU Bits	Entry to replace
0 -> 1 -> 3	000	0
	001	1
0 -> 1 -> 4	010	2
	011	3
0 -> 2 -> 5	100	4
	101	5
0 -> 2 -> 6	110	6
	111	7

Shared Translation Lookaside Buffer

The CVA6 shared TLB is structured as a 2-way associative cache, where the virtual address requiring translation is compared with the set indicated by the virtual page number. The shared TLB is looked up in case of an Instruction TLB (ITLB) or data TLB (DTLB) miss, signaled by these TLBs. If the entry is found in the shared TLB set, the respective TLB, whose translation is being requested, is updated. If the entry is not found in the shared TLB, then the processor has to perform a page table walk. Once the processor obtains a PPN corresponding to the VPN, the shared TLB is updated with this information. If the physical page is not found in the page table, it results in a page fault, which is handled by the operating system. The operating system will then place the corresponding physical page in memory.

The inputs and output signals of the shared TLB are shown in the following two figures.

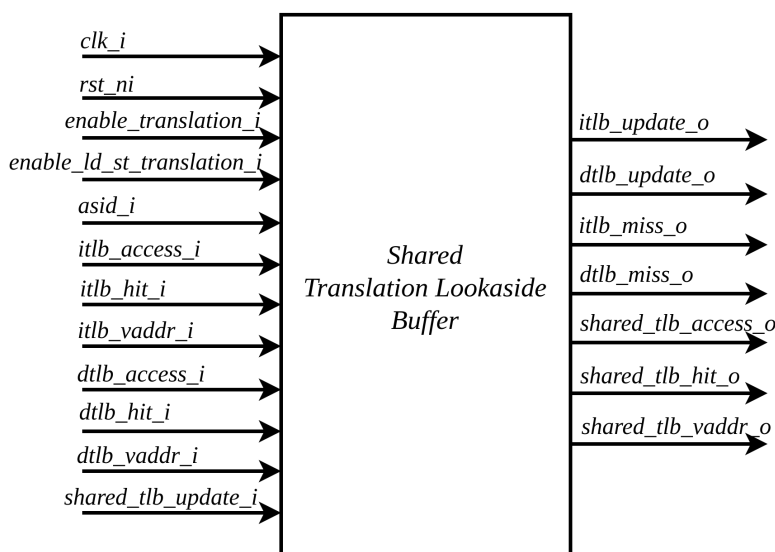


Figure 14: Inputs and outputs of CVA6 shared TLB

3.1 Signal Description

Signal	IO	Connection	Type	Description
<i>clk_i</i>	in	Subsystem	logic	Subsystem Clock
<i>rst_ni</i>	in	Subsystem	logic	Asynchronous reset active low
<i>flush_i</i>	in	Controller	logic	TLB flush request

enable_translation_i	in	CSR Regfile	logic	CSRs indicate to enable SV3
en_ld_st_translation_i	in	CSR Regfile	logic	Enable virtual memory translation for load/stores
asid_i	in	CSR Regfile	logic	ASID for the lookup
itlb_access_i	in	Cache Subsystem	logic	Signal indicating a lookup access in ITLB is being requested.
itlb_hit_i	in	ITLB	logic	Signal indicating a ITLB hit
itlb_vaddr_i	in	Cache Subsystem	logic[31:0]	Virtual address lookup in ITLB
dtlb_access_i	in	Load/Store Unit	logic	Signal indicating a lookup access in DTLB is being requested.
dtlb_hit_i	in	DTLB	logic	Signal indicating a DTLB hit
dtlb_vaddr_i	in	Load/Store Unit	logic[31:0]	Virtual address lookup in DTLB
itlb_update_o	out	ITLB	tlb_update_sv32_t	Tag and content to update ITLB
dtlb_update_o	out	DTLB	tlb_update_sv32_t	Tag and content to update DTLB
itlb_miss_o	out	Performance Counter	logic	Signal indicating ITLB miss
dtlb_miss_o	out	Performance Counter	logic	Signal indicating DTLB miss
shared_tlb_access_o	out	PTW	logic	Signal indicating a lookup access in shared TLB is begin requested
shared_tlb_hit_o	out	PTW	logic	Signal indicating shared TLB hit
shared_tlb_vadd_o	out	PTW	logic[31:0]	Virtual address lookup in shared TLB
itlb_req_o	out	PTW	logic	ITLB Request Output
shared_tlb_update_i	in	PTW	tlb_update_sv32_t	Updated tag and content of shared TLB

3.2 Struct Description

shared_tag_t

Signal	Type	Description
is_4M	logic	Indicates if the shared TLB entry corresponds to a 4MB page.
vpn1	logic[9:0]	Virtual Page Number (VPN) represents the index of PTE in the page table level 1.
vpn0	logic[9:0]	Virtual Page Number (VPN) represents the index of PTE in the page table level 0.
asid	logic	Address Space Identifier (ASID) used to identify different address spaces

3.3 Shared TLB Entry Structure

Shared TLB is 2-way associative, with a depth of 64. A single entry in the set contains the valid bit, tag and the content. The Tag segment stores details such as the virtual page number (VPN1, VPN0), ASID, and page size (is_4M). The Content field contains two physical page numbers (PPN1, PPN0) along with a number of bits which specify various attributes of the physical page.




Figure 15: CVA6 Shared TLB Structure

3.4 Shared TLB Implementation in CVA6

The implementation of a shared TLB in CVA6 is described in the following sections:

- **ITLB and DTLB Miss:** Prepare a shared TLB lookup if the entry is not found in ITLB or DTLB.
- **Tag Comparison:** Look up the provided virtual address in the shared TLB.
- **Update and Flush:** Flush the shared TLB or update it.
- **Replacement Policies:** First non-valid entry and random replacement policy.

3.4.1 ITLB and DTLB Miss

Consider a scenario where an entry is found in the ITLB or DTLB. In this case, there is no need to perform a lookup in the shared TLB since the entry has already been found. Next, there are two scenarios: an ITLB miss or a DTLB miss

To identify an ITLB miss, the following conditions need to be fulfilled:

- Address translation must be enabled.
- There must be an access request to the ITLB.
- The ITLB should indicate an ITLB miss.
- There should be no access request to the DTLB.

During an ITLB miss, access is granted to read the tag and content of the shared TLB from their respective sram. The address for reading the tag and content of the shared TLB entry is calculated using the virtual address for which translation is not found in the ITLB. The ITLB miss is also explicitly indicated by the shared TLB. A request for shared TLB access is initiated.

To identify the DTLB miss, the following conditions need to be fulfilled:

- Address translation for load and stores must be enabled.
- There must be an access request to the DTLB.
- The DTLB should indicate a DTLB miss.

In the case of a DTLB miss, the same logic is employed as described for an ITLB miss.

3.4.2 Tag Comparison

Shared TLB lookup for a hit occurs under the same conditions as described for the TLB modules used as ITLB and DTLB. However, there are some distinctions. In both the ITLB and DTLB, the virtual address requiring translation is compared against all TLB entries. In contrast, the shared TLB only compares the tag and content of the set indicated by the provided virtual page number. **The index of the set is extracted from VPN0 of the requested virtual address.** Given that the shared TLB is 2-way associative, each set contains two entries. Consequently, both of these entries are compared. Below figure illustrates how the set is opted for the lookup.




Figure 16: Set opted for lookup in shared TLB

3.4.3 Update and Flush

Differing from the ITLB and DTLB, a specific virtual address or addressing space cannot be flushed in the shared TLB. When SFENCE.VMA is committed, all entries in the shared TLB are invalidated. **(Cases of SFENCE.VMA should also be added in shared TLB)**

3.4.3.1 Updating Shared TLB

When the Page Table Walker signals a valid update request, the shared TLB is updated by selecting an entry through the replacement policy and marking it as valid. This also triggers the writing of the new tag and content to the respective SRAM.

3.4.4 Replacement Policy Implemented in CVA6 Shared TLB

In CVA6's shared TLB, two replacement policies are employed for replacements based on a specific condition. These replacement policies select the entry within the set indicated by the virtual page number. The two policies are:

- First non-valid encounter replacement policy
- Random replacement policy

First replacement policy failed if all ways are valid. Therefore, a random replacement policy is opted for.

3.4.4.1 First non-valid encounter replacement policy

The module implemented in CVA6 to find the first non-valid entry in the shared TLB is the Leading Zero Counter (LZC). It takes three parameters as input:

1. WIDTH: The width of the input vector.
2. MODE: Mode selection - 0 for trailing zero, 1 for leading zero.
3. CNT WIDTH: Width of the output signal containing the zero count.

The input signal is the vector to be counted, and the output represents the count of trailing/leading zeros. If all bits in the input vector are zero, it will also be indicated.

When initializing the module, the width of the input vector is set to the number of shared TLB ways. The trailing zero counter mode is selected. The vector of valid bits is set as the input vector, but with negation. This is because we want the index of the first non-valid entry, and LZC returns the count of trailing zeros, which actually corresponds to the index of the first occurrence of 1 from the least significant bit (LSB). If there is at least one non-valid entry, that entry is opted for the replacement, and if not then this is signaled by LZC.

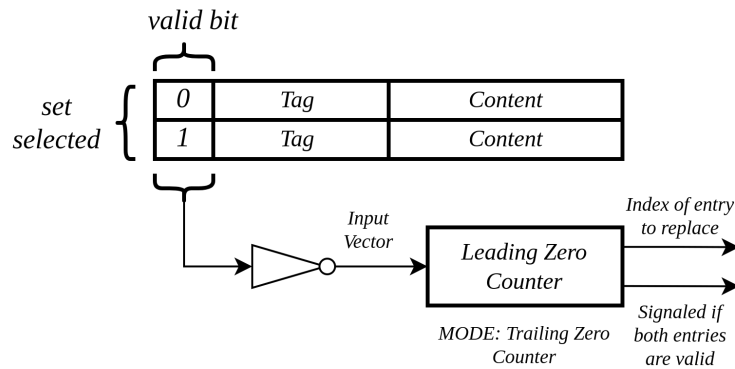


Figure 17: Replacement of First invalid entry.

3.4.4.2 Random replacement policy

If all ways are valid, a random replacement policy is employed for the replacement process. The Linear Feedback Shift Register (LFSR) is utilized to select the replacement entry randomly. LFSR is commonly used in generating sequences of pseudo-random numbers. When the enable signal is active, the current state of the LFSR undergoes a transformation. Specifically, the state is shifted right by one bit, and the result is combined with a predetermined masking pattern. This masking pattern is derived from the predefined “Masks” array, introducing a non-linear behavior to the sequence generation of the LFSR. The masking process involves XOR operations between the shifted state bits and specific pattern bits, contributing to the complexity and unpredictability of the generated sequence.

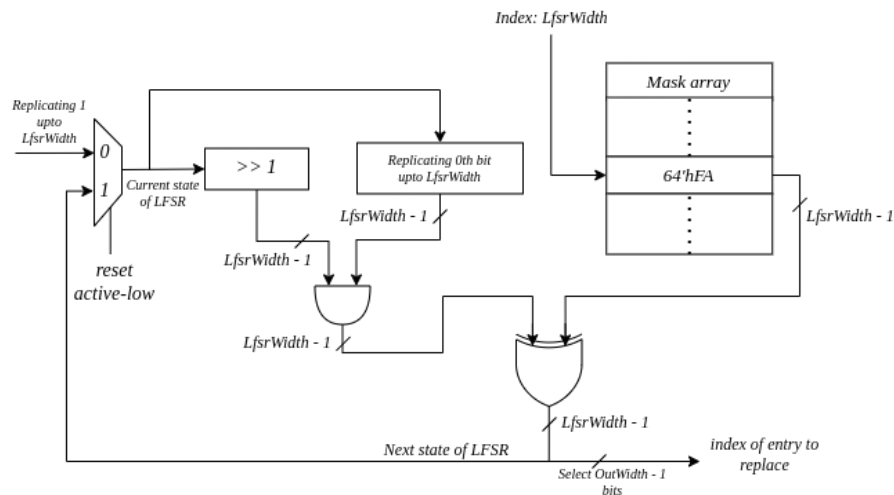


Figure 18: Entry selection for replacement using LFSR

Page Table Walker

The "CVA6 Page Table Walker (PTW) for MMU Sv32" is a hardware module developed for the CV32A6 processor architecture, designed to facilitate the translation of virtual addresses into physical addresses, a crucial task in memory access management.

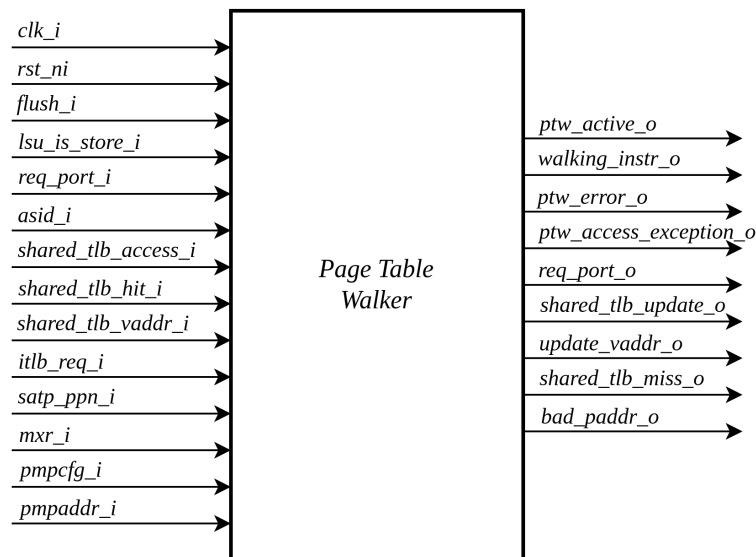


Figure 19: Input and Outputs of Page Table Walker

4.1 Operation of PTW Module

The PTW module operates through various states, each with its specific function, such as handling memory access requests, validating page table entries, and responding to errors.

4.2 Key Features and Capabilities

Key features of this PTW module include support for two levels of page tables (LVL1 and LVL2) in the Sv32 standard, accommodating instruction and data page table walks. It rigorously validates and verifies page table entries (PTEs) to ensure translation accuracy and adherence to access permissions. This module seamlessly integrates with the CV32A6 processor's memory management unit (MMU), which governs memory access control. It also takes into account global mapping, access flags, and privilege

levels during the translation process, ensuring that memory access adheres to the processor's security and privilege settings.

4.3 Exception Handling

In addition to its translation capabilities, the PTW module is equipped to detect and manage errors, including page-fault exceptions and access exceptions, contributing to the robustness of the memory access system. It works harmoniously with physical memory protection (PMP) configurations, a critical aspect of modern processors' memory security. Moreover, the module efficiently processes virtual addresses, generating corresponding physical addresses, all while maintaining speculative translation, a feature essential for preserving processor performance during memory access operations.

4.4 Signal Description

Signal	IO	Connection	Type	Description
clk_i	in	Subsystem	logic	Subsystem Clock
rst_ni	in	Subsystem	logic	Asynchronous reset active low
flush_i	in	Controller	logic	Sfence Committed
ptw_active_o	out	MMU	logic	Output signal indicating whether the Page Table Walker (PTW) is currently active
walking_instr_o	out	MMU	logic	Indicating its a instruction page table walk or not
ptw_error_o	out	MMU	logic	Output signal indicating that an error occurred during PTW operation
ptw_access_exception_o	out	MMU	logic	Output signal indicating that a PMP (Physical Memory Protection) access exception occurred during PTW operation.
lsu_is_store_i	in	Store Unit	logic	Input signal indicating whether the translation was triggered by a store operation.
req_port_i	in	Cache Subsystem	dcache_req_o_t	D Cache Data Requests
req_port_o	out	Cache	dcache_req_u_t	D Cache Data Response

		Subsystem / Perf Counter		
shared_tlb_update_o	out	Shared TLB	tlb_update_sv32_t	Updated tag and content of shared TLB
update_vaddr_o	out	MMU	logic[riscv::VLEN-1:0]	Updated VADDR from shared TLB
asid_i	in	CSR RegFile	logic[ASID_WIDTH-1:0]	ASID for the lookup
shared_tlb_access_i	in	Shared TLB	logic	Access request of shared TLB
shared_tlb_hit_i	in	Shared TLB	logic	Indicate shared TLB hit
shared_tlb_vaddr_i	in	Shared TLB	logic[riscv::VLEN-1:0]	Virtual Address from shared TLB
itlb_req_i	in	Shared TLB	logic	Indicate request to ITLB
satp_ppn_i	in	CSR RegFile	logic[riscv::PPNW-1:0]	PPN of top level page table from SATP register
mxr_i	in	CSR RegFile	logic	Make Executable Readable bit in xSTATUS CSR register
shared_tlb_miss_o	out	OPEN	logic	Indicate a shared TLB miss
pmpcfg_i	in	CSR RegFile	riscv::pmpcfg_t[15:0]	PMP configuration
pmpaddr_i	in	CSR RegFile	logic[15:0][riscv::PLEN-3:0]	PMP Address
bad_paddr_o	out	MMU	logic[riscv::PLEN-1:0]	Bad Physical Address in case of access exception

4.5 Struct Description

dcache_req_i_t

Signal	Type	Description
address_index	logic [DCACHE_INDEX_WIDTH-1:0]	Index of the Dcache Line

address_tag	logic [DCACHE_TAG_ WIDTH-1:0]	Tag of the Dcache Line
data_wdata	riscv::xlen_t	Data to write in the Dcache
data_wuser	logic [DCACHE_USER_ _WIDTH-1:0]	
data_req	logic	Data Request
data_we	logic	Data Write enabled
data_be	logic [(riscv::XLEN/8)-1 :0]	Data Byte enable
data_size	logic [1:0]	
data_id	logic [DCACHE_TID_ WIDTH-1:0]	
kill_req	logic	Kill the D cache request
tag_valid	logic	

dcache_req_o_t

Signal	Type	Description
data_gnt	logic	Grant of data is given in response to the data request
data_rvalid	logic	Indicate that data is valid which is send by D cache
data_rid	logic [DCACHE_TID_ WIDTH-1:0]	
data_rdata	riscv::xlen_t	Data from D cache
data_ruser	logic [DCACHE_USER_ _WIDTH-1:0]	

4.6 PTW State Machine

Page Table Walker is implemented as a finite state machine. It listens to shared TLB for incoming translation requests. If there is a shared TLB miss, it saves the virtual address and starts the page table walk. Page table walker transition between 7 states in CVA6.

- **IDLE:** The initial state where the PTW is awaiting a trigger, often a Shared TLB miss, to initiate a memory access request.
- **WAIT_GRANT:** Request memory access and wait for data grant
- **PTE_LOOKUP:** Once granted access, the PTW examines the valid Page Table Entry (PTE), checking attributes to determine the appropriate course of action.
- **PROPOGATE_ERROR:** If the PTE is invalid, this state handles the propagation of an error, often leading to a page-fault exception due to non-compliance with access conditions
- **PROPOGATE_ACCESS_ERROR:** Propagate access fault if access is not allowed from a PMP perspective
- **WAIT_RVALID:** After processing a PTE, the PTW waits for a valid data signal, indicating that relevant data is ready for further processing.
- **LATENCY:** Introduces a delay to account for synchronization or timing requirements between states.

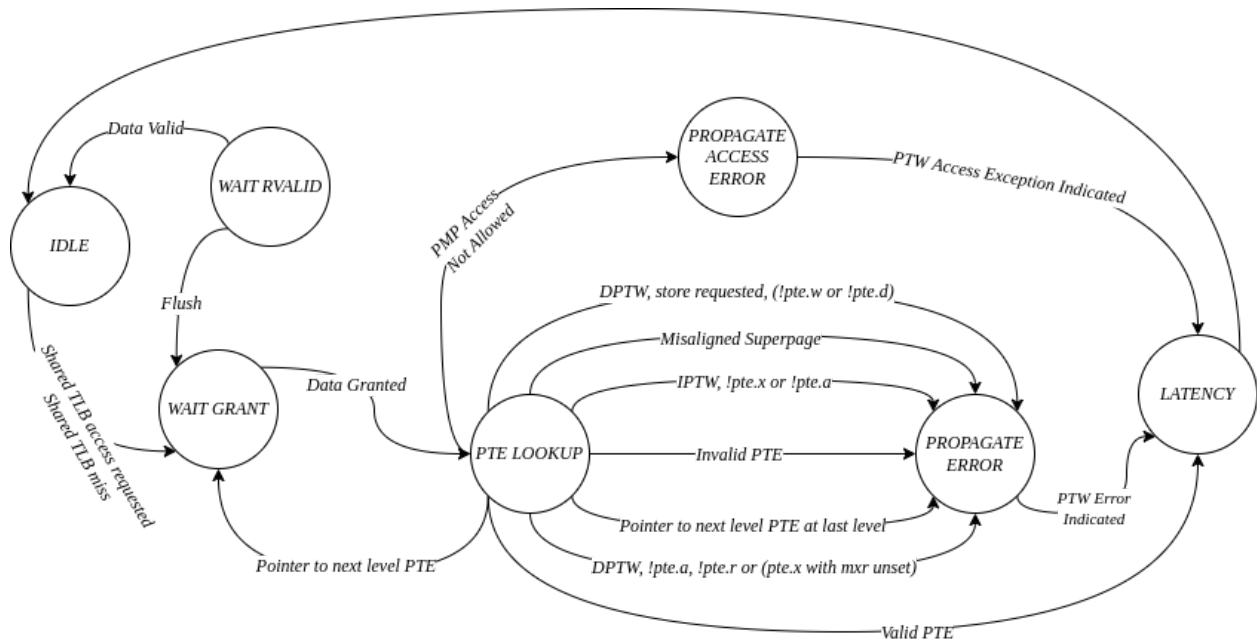


Figure 20: State Machine Diagram of CVA6 PTW

4.6.1 IDLE State

In the **IDLE** state of the Page Table Walker (PTW) finite state machine, the system awaits a trigger to initiate the page table walk process. This trigger is often prompted by a Shared Translation Lookaside Buffer (TLB) miss, indicating that the required translation is not present in the shared TLB cache. The PTW's behavior in this state is explained as follows:

1. The top-most page table is selected for the page table walk. In the case of SV32, which implements a two-level page table, the level 1 page table is chosen.
2. In the **IDLE** state, translations are assumed to be invalid in all addressing spaces.
3. The signal indicating the instruction page table walk is set to 0.
4. A conditional check is performed: if there is a shared TLB access request and the entry is not found in the shared TLB (indicating a shared TLB miss), the following steps are executed:
 - I. The address of the desired Page Table Entry within the level 1 page table is calculated by multiplying the Physical Page Number (PPN) of the level 1 page table from the SATP register by the page size (4kB). This result is then added to the product of the Virtual Page Number (VPN1), and the size of a page table entry(4 bytes).

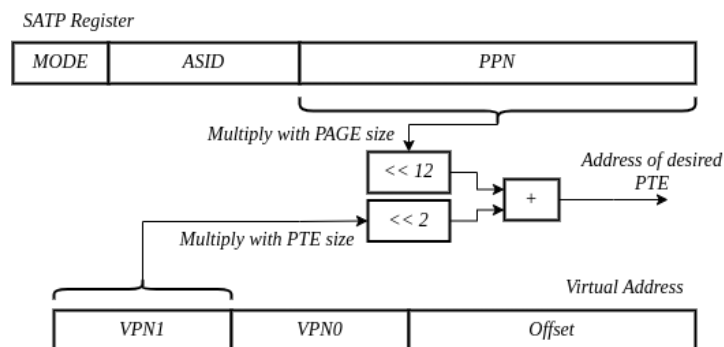


Figure 21: Address of Desired PTE at Level 1

- II. The signal indicating whether it's an instruction page table walk is updated based on the ITLB miss.
- III. The ASID and virtual address are saved for the page table walk.
- IV. A shared TLB miss is indicated.

4.6.2 WAIT GRANT State

In the "WAIT_GRANT" state of the Page Table Walker's finite state machine, a data request is sent to retrieve memory information. It waits for a data grant signal from the Dcache controller, remaining in this state until granted. Once granted, it activates a tag valid signal, marking data validity. The state then transitions to "PTE_LOOKUP" for page table entry lookup.

4.6.3 PTE LOOKUP State

In the "**PTE_LOOKUP**" state of the Page Table Walker (PTW) finite state machine, the PTW performs the actual lookup and evaluation of the page table entry (PTE) based on the virtual address translation. The behavior and operations performed in this state are detailed as follows:

1. The state waits for a valid signal indicating that the data from the memory subsystem, specifically the page table entry, is available for processing.
2. Upon receiving the valid signal, the PTW proceeds with examining the retrieved page table entry to determine its properties and validity.
3. The state checks if the global mapping bit in the PTE is set, and if so, sets the global mapping signal to indicate that the translation applies globally across all address spaces.
4. The state distinguishes between two cases: Invalid PTE and Valid PTE.
 - I. If the valid bit of the PTE is not set, or if the PTE has reserved RWX field encodings, it signifies an Invalid PTE. In such cases, the state transitions to the "**PROPAGATE_ERROR**" state, indicating a page-fault exception due to an invalid translation.




Figure 22: Invalid PTE and reserved RWX encoding leads to page fault

- II. If the PTE is valid, the state advances to the "**LATENCY**" state, indicating a period of processing latency. Additionally, if the "read" flag (pte.r) or the "execute" flag (pte.x) is set, the PTE is considered valid.

5. Within the Valid PTE scenario, the state performs further checks based on whether the translation is intended for instruction fetching or data access:

- I. For instruction page table walk, if the page is not executable (pte.x is not set) or not marked as accessible (pte.a is not set), the state transitions to the "**PROPAGATE_ERROR**" state.

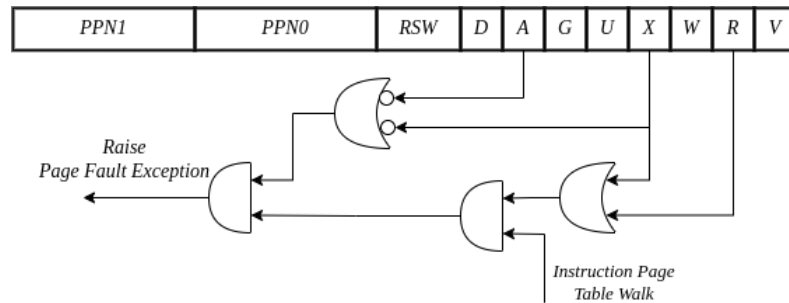


Figure 23: For Instruction Page Table Walk

- II. For data page table walk, the state checks if the page is readable (pte.r is set) or if the page is executable only but made readable by setting the MXR bit in xSTATUS CSR register. If either condition is met, it indicates a valid translation. If not, the state transitions to the "**PROPAGATE_ERROR**" state.

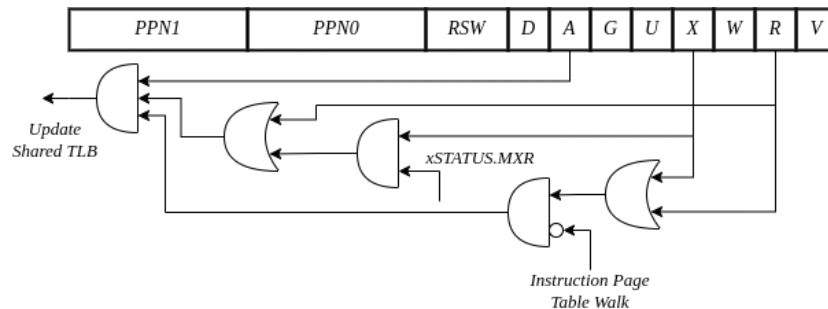


Figure 24: Data Access Page Table Walk

- III. If the access is intended for storing data, additional checks are performed: If the page is not writable (pte.w is not set) or if it is not marked as dirty (pte.d is not set), the state transitions to the "**PROPAGATE_ERROR**" state.

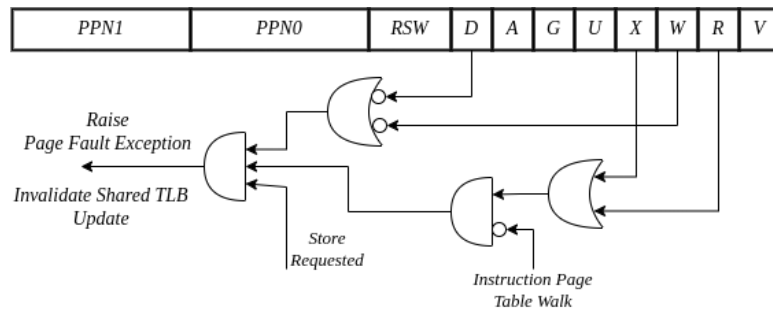


Figure 25: Data Access Page Table Walk, Store requested

6. The state also checks for potential misalignment issues in the translation: If the current page table level is the first level (LVL1) and if the PPN0 of in PTE is not zero, it indicates a misaligned superpage, leading to a transition to the "PROPAGATE_ERROR" state.

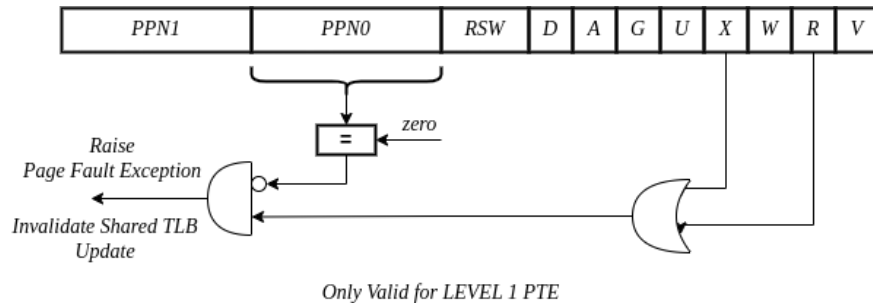


Figure 26: Misaligned Superpage Check

7. If the PTE is valid but the page is neither readable nor executable, the PTW recognizes the PTE as a pointer to the next level of the page table, indicating that additional translation information can be found in the referenced page table at a lower level.

8. If the current page table level is the first level (LVL1), the PTW proceeds to switch to the second level (LVL2) page table, updating the next level pointer and calculating the address for the next page table entry using the Physical Page Number from the PTE and the index of the level 2 page table from virtual address.

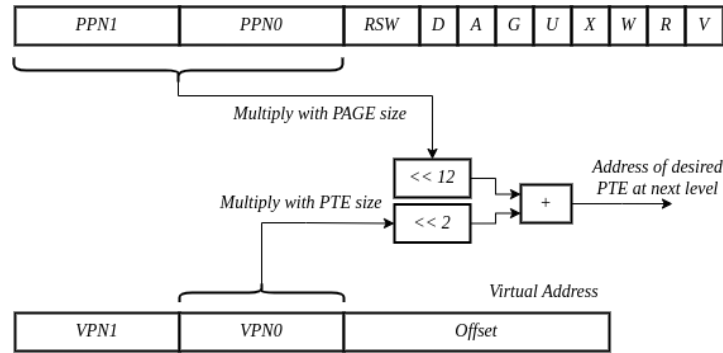


Figure 27: Address of desired PTE at next level of Page Table

9. The state then transitions to the **"WAIT_GRANT"** state, indicating that the PTW is awaiting the grant signal to proceed with requesting the next level page table entry.

10. If the current level is already the second level (LVL2), an error is flagged, and the state transitions to the **"PROPAGATE_ERROR"** state, signifying an unexpected situation where the PTW is already at the last level page table.

11. If the translation access is found to be restricted by the Physical Memory Protection (PMP) settings (*allow_access* is false), the state updates the shared TLB update signal to indicate that the TLB entry should not be updated. Additionally, the saved address for the page table walk is restored to its previous value, and the state transitions to the **"PROPAGATE_ACCESS_ERROR"** state.

12. Lastly, if the data request for the page table entry was granted, the state indicates to the cache subsystem that the tag associated with the data is now valid.

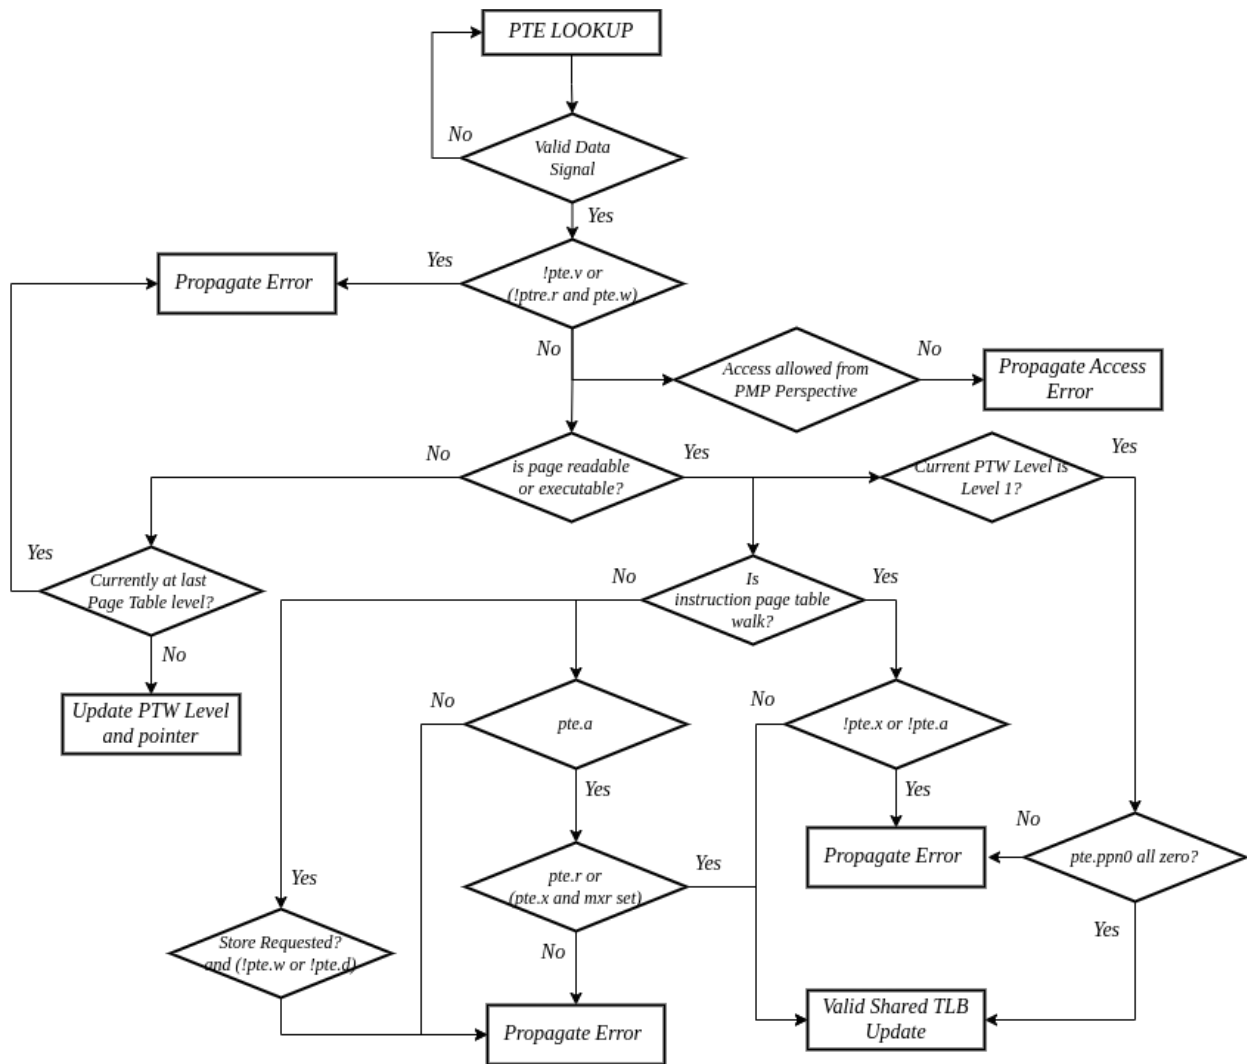


Figure 28: Flow Chart of PTE LOOKUP State

4.6.4 PROPAGATE ERROR State

This state indicates a detected error in the page table walk process, and an error signal is asserted to indicate the Page Table Walker's error condition, triggering a transition to the "LATENCY" state for error signal propagation.

4.6.5 PROPAGATE ACCESS ERROR State

This state indicates a detected access error in the page table walk process, and an access error signal is asserted to indicate the Page Table Walker's access error condition, triggering a transition to the "LATENCY" state for access error signal propagation.

4.6.6 WAIT RVALID State

1. This state waits for the "read valid" signal to be asserted.
2. Once the valid signal is received, the FSM transitions to the IDLE state, indicating that the page table walk process is ready to start again for a new translation.

4.6.7 LATENCY State

The LATENCY state introduces a latency period to allow for necessary system actions or signals to stabilize. After the latency period, the FSM transitions back to the 'IDLE' state, indicating that the system is prepared for a new translation request.

4.7 Flush Scenario

The first step when a flush is triggered is to check whether the Page Table Entry (PTE) lookup process is currently in progress. If the PTW (Page Table Walker) module is indeed in the middle of a PTE lookup operation, the code then proceeds to evaluate a specific aspect of this operation.

4.7.1 Check for Data Validity (rvalid)

Within the PTE lookup operation, it's important to ensure that the data being used for the translation is valid. In other words, the code checks whether the "rvalid" signal (which likely indicates the validity of the data) is not active. If the data is not yet valid, it implies that the PTW module is waiting for the data to become valid before completing the lookup. In such a case, the code takes appropriate action to wait for the data to become valid before proceeding further.

4.7.2 Check for Waiting on Grant

The second condition the code checks for during a flush scenario is whether the PTW module is currently waiting for a "grant." This "grant" signal is typically used to indicate permission or authorization to proceed with an operation. If the PTW module is indeed in a state of waiting for this grant signal, it implies that it requires authorization before continuing its task.

- **Waiting for Grant**

If the PTW module is in a state of waiting for the grant signal, the code ensures that it continues to wait for the grant signal to be asserted before proceeding further.

4.7.3 Return to Idle State if Neither Condition is Met

After evaluating the above two conditions, the code determines whether either of these conditions is true. If neither of these conditions applies, it suggests that the PTW module can return to its idle state, indicating that it can continue normal operations without any dependencies on the flush condition.

